

Bayesian Logistic Regression for Predicting Diabetes Risk Using NHANES 2013–2014 Data

A Capstone Project on Bayesian Applications in Epidemiologic Modeling

Namita Mishra

Autumn Wilcox

2025-11-10

Table of contents

Slides: [slides.html](#) (Edit `slides.qmd`.)

1 Introduction

Diabetes mellitus (DM) remains a major public health challenge, and identifying key risk factors—such as obesity, age, sex, and race/ethnicity—is essential for prevention and targeted intervention. Logistic regression is widely used to estimate associations between such factors and binary outcomes like diabetes diagnosis. However, classical maximum likelihood estimation (MLE) can produce unstable estimates in the presence of missing data, quasi-separation, or small samples. Bayesian logistic regression offers a robust alternative by integrating prior information, regularizing estimates, and quantifying uncertainty more transparently than frequentist approaches.

Bayesian hierarchical models, implemented via Markov Chain Monte Carlo (MCMC), have been successfully applied in predicting patient health status across diseases such as pneumonia, prostate cancer, and mental disorders (?). By representing predictive uncertainty alongside point estimates, Bayesian inference offers a practical advantage in epidemiologic modeling where decisions hinge on probabilistic thresholds. Beyond stability, Bayesian methods support model checking, variable selection, and uncertainty quantification under missingness or imputation frameworks (?; ?).

Recent work has expanded Bayesian applications to disease diagnostics and health risk modeling. For instance, Bayesian approaches have been used to evaluate NHANES diagnostic data (?), to model cardiovascular and metabolic risk (?), and to integrate multiple data modalities such as imaging and laboratory measures (?). Moreover, multiple imputation combined with Bayesian modeling generates robust estimates when data are missing at random (MAR) or not at random (MNAR) (?).

The broader Bayesian literature emphasizes the role of priors and model checking. Weakly informative priors, such as $N(0, 2.5)$ for coefficients, regularize estimation and reduce variance in small samples (?; ?). Tutorials using R packages like `brms` and `blavaan` illustrate how MCMC enables posterior inference and empirical Bayes analysis (?).

Beyond standard generalized linear models, Bayesian nonparametric regression flexibly captures nonlinearity and zero inflation common in health data (?). Bayesian Additive Regression Trees (BART) improve variable selection in mixed-type data (?), while state-space and dynamic Bayesian models incorporate time-varying biomarkers for longitudinal prediction (?). Bayesian model averaging (BMA) further addresses model uncertainty by weighting across multiple specifications (?). Together, these approaches demonstrate the versatility and growing importance of Bayesian inference in clinical and epidemiologic modeling.

The objective of this project is to evaluate whether Bayesian inference provides more stable and interpretable estimates of diabetes risk than frequentist and imputation-based approaches, particularly when data complexity or separation challenges arise. Agreement across modeling frameworks supports the robustness of these associations and highlights the interpretability and uncertainty quantification advantages offered by Bayesian analysis in population health modeling (?).

1.1 Aims

The present study employs Bayesian logistic regression to predict diabetes status and examine the relationships between diabetes and key predictors, including body mass index (BMI), age (20 years), sex, and race. Using retrospective data from the 2013–2014 NHANES survey, the analysis accounts for the study’s complex sampling design, which involves stratification, clustering, and the oversampling of specific subpopulations rather than simple random sampling. The Bayesian framework is applied to address common analytical challenges such as missing data, complete case bias, and data separation, thereby improving the robustness and reliability of inference compared to traditional logistic regression methods.

2 Method

2.1 Bayesian Logistic Regression

The Bayesian framework integrates prior knowledge with observed data to generate posterior distributions, allowing parameters to be interpreted directly in probabilistic terms.

Unlike traditional frequentist approaches that yield single-point estimates and p-values, Bayesian methods represent parameters as random variables with full probability distributions.

This provides greater flexibility, incorporates parameter uncertainty, and produces credible intervals that directly quantify the probability that a parameter lies within a given range.

2.2 Model Structure

Bayesian logistic regression models the log-odds of a binary outcome as a linear combination of predictors:

$$\text{logit}(P(Y = 1)) = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \cdots + \beta_k X_k$$

where

- $P(Y = 1)$ is the probability of the event of interest,
- β_0 is the intercept (log-odds when all predictors are zero), and
- β_j represents the effect of predictor X_j on the log-odds of the outcome, holding other predictors constant.

In the Bayesian framework, model parameters (β) are treated as random variables and assigned prior distributions that reflect existing knowledge or plausible ranges before observing the data. After incorporating the observed evidence, the priors are updated through Bayes' theorem (?; ?):

$$\text{Posterior} \propto \text{Likelihood} \times \text{Prior}$$

- **Likelihood:** represents the probability of the observed data given the model parameters—it captures how well different parameter values explain the data.
- **Prior:** expresses beliefs or existing information about the parameters before observing the data.
- **Posterior:** combines both, representing the updated distribution of parameter values after accounting for the data.

This formulation allows uncertainty to propagate naturally through the model, producing posterior distributions for each coefficient that can be directly interpreted as probabilities.

2.3 Prior Specification

Weakly informative priors were used to regularize estimation without imposing strong assumptions:

- **Regression Coefficients:** $N(0, 2.5)$, providing gentle regularization while allowing substantial variation in plausible effects (?; ?).
- **Intercept:** Student's t-distribution prior, $t(3, 0, 10)$ (?; ?), which has
 - 3 degrees of freedom (heavy tails to allow occasional large effects),
 - mean 0 (no bias toward positive or negative effects), and
 - scale 10 (broad range of possible values).

Such priors help stabilize estimation in the presence of multicollinearity, limited sample size, or potential outliers.

2.4 Posterior Predictions

Posterior distributions of regression coefficients were used to estimate the probability of the outcome for given predictor values. This allows statements such as: > Given the predictors, the probability of the outcome lies between X% and Y%.

Posterior predictions account for two key sources of uncertainty:

1. **Parameter Uncertainty:** Variability in estimated model coefficients.
2. **Predictive Uncertainty:** Variability in possible future outcomes given those parameters.

In Bayesian analysis, all unknown quantities—coefficients, means, variances, or probabilities—are treated as random variables described by their posterior distributions.

2.5 Model Evaluation and Diagnostics

Model quality and convergence were assessed using standard Bayesian diagnostics:

- **Posterior Sampling:** Conducted via Markov Chain Monte Carlo (MCMC) using the No-U-Turn Sampler (NUTS), a variant of Hamiltonian Monte Carlo (HMC) (?). Four chains were run with sufficient warm-up iterations to ensure convergence.
- **Convergence Metrics:** The potential scale reduction factor (\hat{R}) and effective sample size (ESS) were used to verify stability and mixing across chains.
- **Autocorrelation Checks:** Ensured independence between successive draws.
- **Posterior Predictive Checks (PPCs):** Compared simulated outcomes to observed data to evaluate fit.
- **Bayesian R^2 :** Quantified the proportion of variance explained by predictors, incorporating posterior uncertainty.

2.6 Advantages of Bayesian Logistic Regression

- **Uncertainty Quantification:** Produces full posterior distributions instead of single estimates.
- **Credible Intervals:** Provide the range within which a parameter lies with a specified probability (e.g., 95%).
- **Flexible Priors:** Allow integration of expert knowledge or findings from prior studies.
- **Probabilistic Predictions:** Posterior predictive distributions yield direct probabilities for new or future observations.
- **Model Evaluation:** PPCs assess how well simulated outcomes reproduce observed data.

3 Analysis and Results

3.1 Data Preparation

This study used publicly available 2013–2014 NHANES data published by the CDC’s National Center for Health Statistics (?). Three component files were utilized: DEMO_H (demographics), BMX_H (body measures), and DIQ_H (diabetes questionnaire). Each file was imported in .XPT format using the **haven** package in **R**, and merged using the unique participant identifier SEQN to create a single adult analytic dataset (age 20 years).

All variables were coerced to consistent numeric or factor types prior to merging to ensure atomic columns suitable for survey-weighted analysis and modeling. The use of SEQN preserved respondent integrity across datasets and ensured accurate record linkage. This preprocessing step standardized variable formats and minimized inconsistencies between files.

Data wrangling, cleaning, and merging were performed in **R** using a combination of base functions and tidyverse packages. Bayesian logistic regression modeling was later implemented using the **brms**

interface to **Stan**, allowing probabilistic inference within a reproducible workflow that accommodated the NHANES complex survey design and missing data considerations.

3.1.1 Data Import and Merging

```
merged_data <- readRDS("data/merged_2013_2014.rds")  
  
merged_n <- nrow(merged_data)
```

The merged dataset contains 10,175 participants. It integrates demographic, examination, and diabetes questionnaire data. We then restrict the sample to adults (age ≥ 20) to define the analytic cohort used in subsequent analyses. A small proportion of records have missing values in BMI and diabetes status, which will be addressed later through multiple imputation.

Table 1: Preview of merged NHANES 2013–2014 dataset limited to analysis variables (source columns only).

RIDAGEYR	BMXBMI	RIAGENDR	RIDRETH1	DIQ010
69	26.7	1	4	1
54	28.6	1	3	1
72	28.9	1	3	1
9	17.1	1	3	2
73	19.7	2	3	2
56	41.7	1	1	2
0	NA	1	3	NA
61	35.7	2	3	2
42	NA	1	2	2
56	26.5	2	3	2

3.1.2 Variable Definitions

- **Response Variable:**

`diabetes_dx` (binary) represents a Type 2 diabetes diagnosis, excluding gestational diabetes. It was derived from `DIQ010` (“Doctor told you have diabetes”), while `DIQ050` (insulin use) was excluded to prevent treatment-related confounding.

- **Predictor Variables:**

- `BMXBMI` – Body Mass Index (kg/m^2), treated as continuous and categorized into six BMI classes (`bmi_cat`).
- `RIDAGEYR` – Age (continuous, 20–80 years)
- `RIAGENDR` – Sex (factor, two levels)

- RIDRETH1 – Ethnicity (factor, five levels)

```
# -----
# Variable descriptions (with `code` formatting for names)
# -----
var_tbl <- tribble(
  ~Variable,      ~Description,
  "``diabetes_dx``", "Type 2 diabetes diagnosis (1 = Yes, 0 = No) derived from `DIQ010`; gestation",
  "``age``",       "Age in years.",
  "``bmi``",        "Body Mass Index (kg/m^2) computed from measured height and weight.",
  "``bmi_cat``",    "BMI categories: Underweight, Normal, Overweight, Obesity I-III (`Normal` is",
  "``sex``",        "Sex of participant (`Male`, `Female`).",
  "``race``",      "race/Ethnicity collapsed to four levels: White, Black, Hispanic, Other.", "Categori",
  "``WTMEC2YR``",   "Examination sample weight for Mobile Examination Center participants.",
  "``SDMVPSU``",    "Primary Sampling Unit used for variance estimation in the complex survey de",
  "``SDMVSTRA``",   "Stratum identifier used to define strata for the complex survey design.",
  "``age_c``",       "Centered and standardized age (z-score).",
  "``bmi_c``",       "Centered and standardized BMI (z-score).",
)

kbl(
  var_tbl,
  caption = "Variable Descriptions: Adult Analytic Dataset",
  align = c("l","l","l","l"),
  escape = FALSE
) %>%
  kable_styling(full_width = FALSE, position = "center", bootstrap_options = c("striped","hover")
  group_rows("Analysis variables", 1, 6) %>%                      # <-- updated range (now 6 analysis
  group_rows("Survey design variables", 7, 9) %>%
  group_rows("Derived variables", 10, 11)
```

3.1.3 Study Design and Survey-Weighted Analysis

The National Health and Nutrition Examination Survey (NHANES) employs a complex, multistage probability sampling design with stratification, clustering, and oversampling of specific demographic groups (for example, racial/ethnic minorities and older adults) to produce nationally representative estimates of the U.S. population.

Survey design variables — primary sampling units (**SDMVPSU**), strata (**SDMVSTRA**), and examination sample weights (**WTMEC2YR**) — were retained to account for this complex design. These variables were applied to adjust for unequal probabilities of selection, nonresponse, and oversampling, ensuring valid standard errors, unbiased prevalence estimates, and generalizable population-level inference.

A survey-weighted logistic regression model was used to evaluate the association between diabetes status (**diabetes_dx**, binary outcome) and key predictors: body mass index (**bmi**), age (**age**), sex (**sex**), and race/ethnicity (**race**). Diabetes was defined using **DIQ010** (“Doctor told you

Table 2: Variable Descriptions: Adult Analytic Dataset

Variable	Description
Analysis variables	
‘diabetes _{dx} ’	Type 2 diabetes diagnosis (1 = Yes, 0 = No) derived from ‘DIQ010’; gestational diabetes excluded.
‘age’	Age in years.
‘bmi’	Body Mass Index (kg/m^2) computed from measured height and weight.
‘bmi _{cat} ’	BMI categories: Underweight, Normal, Overweight, Obesity I–III (‘Normal’ is reference in model).
‘sex’	Sex of participant (‘Male’, ‘Female’).
‘race’	race/Ethnicity collapsed to four levels: White, Black, Hispanic, Other.
Survey design variables	
‘WTMEC2YR’	Examination sample weight for Mobile Examination Center participants.
‘SDMVPSU’	Primary Sampling Unit used for variance estimation in the complex survey design.
‘SDMVSTRA’	Stratum identifier used to define strata for the complex survey design.
Derived variables	
‘age _c ’	Centered and standardized age (z-score).
‘bmi _c ’	Centered and standardized BMI (z-score).

have diabetes”) and coded as 0/1, with DIQ050 (insulin use) excluded to avoid treatment-related confounding.

Covariates included:

- `age` (continuous; centered as `age_c`, categorized 20–80 years)
- `bmi` (continuous; centered as `bmi_c`, and categorized by BMI class `bmi_cat`)
- `sex` (male, female)
- `race` (four ethnicity levels: White, Black, Hispanic, Other)

This approach accounts for NHANES’ complex sampling design, producing unbiased parameter estimates and valid inference for U.S. adults.

Step	Description
Weighting	Used the <code>survey</code> package to calculate weighted means for key variables (e.g., age and diabetes status) and to estimate design effects and effective sample size for the complex survey design.
Standardization	Centered and standardized BMI and age (<code>bmi_c</code> , <code>age_c</code>) for use in regression models.
Age Categorization	Not implemented in the analytic dataset (continuous <code>age</code> retained). Reference retained for potential descriptive grouping (20–<30, 30–<40, 40–<50, 50–<60, 60–<70, 70–80).
BMI Categorization	Recoded as: <18.5 (Underweight), 18.5–<25 (Normal), 25–<30 (Overweight), 30–<35 (Obesity I), 35–<40 (Obesity II), 40 (Obesity III).
Ethnicity Recoding	<code>RIDRETH1</code> recoded as: 1 = Mexican American, 2 = Other Hispanic, 3 = Non-Hispanic White, 4 = Non-Hispanic Black, 5 = Other/Multi; then <code>NH_White</code> set as the reference level (five analytical levels retained).

Step	Description
Special Codes	Transformed nonresponse codes (e.g., 3, 7, 9) to <code>NA</code> . These missing codes were evaluated for potential nonrandom patterns (MAR/MNAR).
Missing Data	Retained and visualized missing values (primarily in BMI and diabetes status) to assess their pattern and informativeness before multiple imputation.
Final Dataset	Created the cleaned analytic dataset (<code>adult</code>) using <i>Non-Hispanic White</i> and <i>Male</i> as reference groups for modeling, preserving NHANES survey design variables (<code>WTMEC2YR</code> , <code>SDMVPSU</code> , <code>SDMVSTRA</code>).

3.1.4 Adult Cohort Definition

```
# NHANES survey design object for the adult analytic cohort

nhanes_design_adult <- survey::svydesign(
  id      = ~SDMVPSU,
  strata  = ~SDMVSTRA,
  weights = ~WTMEC2YR,
  nest    = TRUE,
  data    = adult
)

# Quick weighted checks

survey::svymean(~age, nhanes_design_adult, na.rm = TRUE)
```

```
mean      SE
age 47.496 0.3805
```

```
survey::svymean(~diabetes_dx, nhanes_design_adult, na.rm = TRUE)
```

```
mean      SE
diabetes_dx 0.089016 0.0048
```

```
# Design effect and effective sample size for `diabetes_dx`

v_hat <- as.numeric(survey::svyvar(~diabetes_dx, nhanes_design_adult, na.rm = TRUE))
p_hat <- mean(adult$diabetes_dx, na.rm = TRUE)
n_obs <- nrow(adult)
v_srs <- p_hat * (1 - p_hat) / n_obs
deff  <- v_hat / v_srs
```

```

n_total <- sum(weights(nhanes_design_adult), na.rm = TRUE)
ess      <- as.numeric(n_total / deff)

cat("Design effect for diabetes_dx:", round(deff, 2), "\n")

```

Design effect for diabetes_dx: 4759.91

```

cat("Effective sample size for diabetes_dx:", round(ess), "\n")

```

Effective sample size for diabetes_dx: 48142

Descriptive statistics for continuous and categorical variables are presented below.

```

# Keep only analytic variables for Table 1
tbl1_dat <- adult %>%
  transmute(
    age,
    bmi,
    bmi_cat,
    sex,
    race,
    diabetes_dx = factor(diabetes_dx, levels = c(0, 1), labels = c("No", "Yes"))
  )

# Continuous summaries: N, missing, mean, sd, min, max
cont_vars <- c("age", "bmi")

cont_sum <- tbl1_dat %>%
  select(all_of(cont_vars)) %>%
  pivot_longer(everything(), names_to = "Variable", values_to = "value") %>%
  group_by(Variable) %>%
  summarise(
    N      = sum(!is.na(value)),
    Missing = sum(is.na(value)),
    Mean   = round(mean(value, na.rm = TRUE), 2),
    SD     = round(sd(value, na.rm = TRUE), 2),
    Min    = round(min(value, na.rm = TRUE), 1),
    Max    = round(max(value, na.rm = TRUE), 1),
    .groups = "drop"
  )

# Categorical summaries: counts and percents
cat_vars <- c("sex", "race", "diabetes_dx", "bmi_cat")

cat_sum <- tbl1_dat %>%

```

```

mutate(across(all_of(cat_vars),
             ~ forcats::fct_explicit_na(as.factor(.x), na_level = "(Missing)")) %>%
select(all_of(cat_vars)) %>%
pivot_longer(everything(), names_to = "Variable", values_to = "Level") %>%
count(Variable, Level, name = "n") %>%
group_by(Variable) %>%
mutate(pct = round(100 * n / sum(n), 1)) %>%
ungroup() %>%
arrange(Variable, desc(n))

# Render tables
kable(cont_sum,
      caption = "Table 1a. Continuous variables (age, BMI): N, missing, mean (SD), range.") %>%
kable_styling(full_width = FALSE)

```

Table 4: Table 1a. Continuous variables (age, BMI): N, missing, mean (SD), range.

Variable	N	Missing	Mean	SD	Min	Max
age	5769	0	49.11	17.56	20.0	80.0
bmi	5520	249	29.10	7.15	14.1	82.9

```

kable(cat_sum,
      caption = "Table 1b. Categorical variables (sex, race, diabetes_dx, bmi_cat): counts and percentages." %>%
kable_styling(full_width = FALSE)

```

Table 5: Table 1b. Categorical variables (sex, race, diabetes_{dx}, bmi_{cat}) : counts and percentages.

Variable	Level	n	pct
bmi_cat	25-<30	1768	30.6
bmi_cat	18.5-<25	1579	27.4
bmi_cat	30-<35	1145	19.8
bmi_cat	35-<40	519	9.0
bmi_cat	40	419	7.3
bmi_cat	(Missing)	249	4.3
bmi_cat	<18.5	90	1.6
diabetes_dx	No	4974	86.2
diabetes_dx	Yes	618	10.7
diabetes_dx	(Missing)	177	3.1
race	NH White	2472	42.8
race	NH Black	1177	20.4
race	Other/Multi	845	14.6
race	Mexican American	767	13.3
race	Other Hispanic	508	8.8

sex	Female	3011	52.2
sex	Male	2758	47.8

Table 1a and 1b summarize the analytic variables included in subsequent models. Mean age and BMI values indicate an adult cohort spanning a wide range of body composition, while categorical summaries confirm balanced sex representation and sufficient sample sizes across race/ethnicity categories. These variables were standardized and used as predictors in all modeling frameworks (analytic cohort N = 5,769 adults 20 years).

```
adult_n <- nrow(adult)
```

Table 6: Excerpt of the cleaned NHANES 2013–2014 adult cohort (age 20; N = 5,769) with derived and standardized variables.

SDMVPSU	SDMVSTRA	WTMEC2YR	diabetes_dx	bmi	age	sex	race	DIQ050	age_c	bmi_c	bmi_cat
1	112	13481.04	1	26.7	69	Male	NH Black	1	1.1324183	-	25–0.3358861<30
1	108	24471.77	1	28.6	54	Male	NH White	1	0.2783598	-	25–0.0702810<30
1	109	57193.29	1	28.9	72	Male	NH White	1	1.3032300	-	25–0.0283434<30
2	116	65541.87	0	19.7	73	Female	NH White	2	1.3601672	-	18.5–1.3144311<25
1	111	25344.99	0	41.7	56	Male	Mexican American	2	0.39223431.760996140		
1	114	61758.65	0	35.7	61	Female	NH White	2	0.67692040.922243235–<40		

As shown in Table ??, the analytic adult cohort (N = 5,769) includes standardized variables for age and BMI (`age_c`, `bmi_c`), categorical indicators for sex and race/ethnicity (`race`), and a binary doctor-diagnosed diabetes variable (`diabetes_dx`).

```
# Textual structure and preview
str(adult)
```

```
'data.frame': 5769 obs. of 12 variables:
 $ SDMVPSU   : num  1 1 1 2 1 1 2 1 2 2 ...
 $ SDMVSTRA  : num  112 108 109 116 111 114 106 112 112 113 ...
 $ WTMEC2YR  : num  13481 24472 57193 65542 25345 ...
 $ diabetes_dx: num  1 1 1 0 0 0 0 0 0 0 ...
 $ bmi        : num  26.7 28.6 28.9 19.7 41.7 35.7 NA 26.5 22 20.3 ...
 $ age        : num  69 54 72 73 56 61 42 56 65 26 ...
 $ sex        : Factor w/ 2 levels "Male","Female": 1 1 1 2 1 2 1 2 1 2 ...
 $ race       : Factor w/ 5 levels "NH White","Mexican American",...: 4 1 1 1 2 1 3 1 1 1 ...
 $ DIQ050     : num  1 1 1 2 2 2 2 2 2 2 ...
```

```
$ age_c      : num  1.132 0.278 1.303 1.36 0.392 ...
$ bmi_c      : num  -0.3359 -0.0703 -0.0283 -1.3144 1.761 ...
$ bmi_cat    : Factor w/ 6 levels "<18.5","18.5-<25",...: 3 3 3 2 6 5 NA 3 2 2 ...
```

```
# Visual structure and type overview
plot_intro(adult, title = "Adult Dataset: Variable Types and Completeness")
```




Figure 1: The visual overview indicates that 75% of variables are continuous and 25% are categorical, with no completely missing columns. Approximately 92.7% of rows are fully complete, and only 1.3% of observations contain missing values, suggesting minimal data loss prior to imputation.

3.1.5 Missing Data Summary

```
# Visualize missing data pattern
plot_missing(adult, title = "Missing Data Pattern (Adult Dataset)")
```




Figure 2: Missing data were minimal across analytic variables. BMI-related fields (bmi, bmi_c, bmi_cat) showed ~4.3% missingness, and diabetes_dx showed ~3.1%. All demographic and survey design variables were complete, indicating that missingness was limited to health-related measures and appropriate for multiple imputation.

```
# Summarize missingness for key analysis variables

miss_tbl <- tibble::tibble(
  Variable = c("bmi", "diabetes_dx"),
  Missing_n = c(sum(is.na(adult_eda$bmi)), sum(is.na(adult_eda$diabetes_dx))),
  Missing_pct = round(c(mean(is.na(adult_eda$bmi)), mean(is.na(adult_eda$diabetes_dx))) * 100, 1)
)

knitr::kable(
  miss_tbl,
  caption = "Missingness for Key Analysis Variables."
)
```

Table 7: Missingness for Key Analysis Variables.

Variable	Missing_n	Missing_pct
bmi	249	4.3
diabetes_dx	177	3.1

Overall missingness was (~7.3%). Gaps were concentrated in bmi ($n = 249$) and diabetes_dx ($n = 177$), while demographic and design variables were complete. This limited pattern of missingness

is consistent with a Missing At Random (MAR) mechanism and likely reflects reduced participation in the physical examination component among certain adults.

3.1.6 Exploratory Data Analysis

Following the missing data assessment, exploratory analyses were conducted to describe the adult analytic cohort and visualize distributions across key demographic and health variables. The goal was to examine univariate patterns and bivariate relationships relevant to diabetes prevalence prior to modeling.

The adult analytic cohort was broadly representative of the U.S. population, with a majority identifying as Non-Hispanic White. Age and BMI distributions were right-skewed, with most participants classified as overweight or obese. Visual exploration revealed a clear positive association between age, BMI, and diabetes prevalence. Non-Hispanic Black and Hispanic participants exhibited higher diabetes prevalence compared with Non-Hispanic Whites.

Approximately 25% of variables were categorical (e.g., sex, race, diabetes_dx) and 75% were continuous (e.g., age, bmi, age_c, bmi_c), indicating that the dataset primarily comprised measured numeric values. About 93% of observations contained complete information across all predictors and outcomes, reflecting high data quality.

Adult age ranged from 20 to 80 years, with peak representation between 30 and 50 years and a slight right skew toward older ages. BMI was concentrated in the overweight and obese ranges, and Female participants were slightly overrepresented relative to Male participants.

```
# Age distribution (analytic adult)
ggplot(adult, aes(x = age)) +
  geom_histogram(binwidth = 5, color = "white") +
  labs(title = "Distribution of Age (20 years)", x = "Age (years)", y = "Count") +
  theme_minimal()
```




Figure 3: Distribution of age among adults aged 20 years. The sample spans 20–80 years, with peak representation between 30 and 50 years and a gradual decline in older age groups, reflecting a balanced adult cohort suitable for regression modeling.

```
# Diabetes outcome distribution
ggplot(adult, aes(x = factor(diabetes_dx, levels = c(0,1), labels = c("No","Yes")))) +
  geom_bar() +
  labs(title = "Diabetes Outcome Distribution (20 years)", x = "Diabetes (No/Yes)", y = "Count")
  theme_minimal()
```




Figure 4: Distribution of diabetes outcomes among adults aged 20 years. Most participants reported no diabetes diagnosis (**No**), while approximately 11% had diabetes (**Yes**) and 3% had missing responses, reflecting expected population prevalence and limited outcome missingness.

```
# BMI category distribution
ggplot(adult, aes(x = bmi_cat)) +
  geom_bar(color = "white", fill = "skyblue") +
  labs(title = "Distribution of BMI Categories (20 years)", x = "BMI Category", y = "Count") +
  theme_minimal()
```




Figure 5: Distribution of BMI categories among adults aged 20 years. The majority of participants fell within the overweight ($25\text{--}<30$) and obese (30) ranges, with fewer individuals classified as underweight (<18.5). This distribution aligns with national trends in adult body composition, supporting the dataset's representativeness for metabolic health analyses.

```
# BMI by diabetes outcome (boxplot)
# (You can't use boxplot with categorical y, so revert to numeric BMI here)
ggplot(adult, aes(x = factor(diabetes_dx, levels = c(0,1), labels = c("No","Yes")), y = bmi)) +
  geom_boxplot(fill = "lightblue") +
  labs(title = "BMI by Diabetes Diagnosis (20 years)", x = "Diabetes (No/Yes)", y = "BMI (numerical)")
  theme_minimal()
```




Figure 6: Distribution of BMI by diabetes diagnosis among adults aged 20 years. Participants with diabetes (**Yes**) show a higher median BMI and greater variability compared to those without diabetes (**No**), supporting the established positive association between obesity and diabetes risk.

```
# Diabetes by race (dodged bars)
ggplot(adult, aes(x = race, fill = factor(diabetes_dx, levels = c(0,1), labels = c("No","Yes")))
  geom_bar(position = "dodge") +
  labs(title = "Diabetes Diagnosis by race/Ethnicity (20 years)",
       x = "race/Ethnicity (race)", y = "Count", fill = "Diabetes") +
  theme_minimal() +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```




Figure 7: Diabetes diagnosis by race/ethnicity among adults aged 20 years. Non-Hispanic Black and Hispanic participants show higher proportions of diabetes diagnoses compared with Non-Hispanic White participants, reflecting known disparities in diabetes prevalence across racial and ethnic groups.

3.2 Modeling Frameworks

Three modeling frameworks were compared using identical predictors—standardized age, BMI, sex, and race—and the binary outcome `diabetes_dx`:

- (1) survey-weighted logistic regression to account for the NHANES complex sampling design,
- (2) multiple imputation (MICE) to handle missing BMI values, and
- (3) Bayesian logistic regression with weakly informative priors to quantify parameter uncertainty.

3.2.1 Survey-Weighted Logistic Regression (Design-Based MLE)

The NHANES 2013–2014 data use a complex, multistage probability design involving strata (`SDMVSTRA`), primary sampling units (PSUs; `SDMVPSU`), and examination weights (`WTMEC2YR`) to ensure nationally representative estimates (?).

Estimates are population-weighted using NHANES survey design variables (`WTMEC2YR`, `SDMVSTRA`, `SDMVPSU`). Odds ratios are reported per one standard deviation (1 SD) increase in age and BMI, with reference groups Male and White.

```

adult_clean <- adult %>%
  dplyr::mutate(
    sex    = forcats::fct_drop(sex),
    race   = forcats::fct_drop(race),
    age_c  = as.numeric(age_c),
    bmi_c  = as.numeric(bmi_c)
  ) %>%
  dplyr::filter(
    !is.na(diabetes_dx),
    !is.na(age_c),
    !is.na(bmi_c),
    !is.na(sex),
    !is.na(race)
  )

```

Below is a structure of the analytic dataset used for regression modeling, showing variable names, types, and sample values (N = 5,349).

```
str(adult_clean[, c("diabetes_dx", "sex", "race", "age_c", "bmi_c")])
```

```
'data.frame': 5349 obs. of 5 variables:
$ diabetes_dx: num 1 1 1 0 0 0 0 0 0 1 ...
$ sex          : Factor w/ 2 levels "Male","Female": 1 1 1 2 1 2 2 1 2 1 ...
$ race         : Factor w/ 5 levels "NH White","Mexican American",...: 4 1 1 1 2 1 1 1 1 ...
$ age_c        : num 1.132 0.278 1.303 1.36 0.392 ...
$ bmi_c        : num -0.3359 -0.0703 -0.0283 -1.3144 1.761 ...
```

```
knitr::kable(
  table(adult_clean$sex)
)
```

Table 8: Distribution of participants by sex (Male = 2,551; Female = 2,798) in the analytic cohort.

	Var1	Freq
Male	2551	
Female	2798	

```
knitr::kable(
  table(adult_clean$race)
)
```

Table 9: Race/ethnicity composition of the analytic cohort, with most participants identifying as Non-Hispanic White (n = 2,293) and Non-Hispanic Black (n = 1,101).

	Var1	Freq
NH White		2293
Mexican American		713
Other Hispanic		470
NH Black		1101
Other/Multi		772

```
knitr::kable(
  table(adult_clean$diabetes_dx)
)
```

Table 10: Observed diabetes prevalence (binary outcome variable `diabetes_dx`), with 597 diagnosed cases (1 = Yes) and 4,752 non-diabetic participants (0 = No).

	Var1	Freq
0		4752
1		597

```
options(survey.lonely.psu = "adjust")

nhanes_design_adult <- survey::svydesign(
  id      = ~SDMVPSU,
  strata  = ~SDMVSTRA,
  weights = ~WTMEC2YR,
  nest    = TRUE,
  data    = adult_clean
)

svy_fit <- survey::svyglm(
  diabetes_dx ~ age_c + bmi_c + sex + race,
  design = nhanes_design_adult,
  family = quasibinomial()
)

svy_or <- broom::tidy(svy_fit, conf.int = TRUE) %>%
  dplyr::mutate(
    OR   = exp(estimate),
    LCL = exp(conf.low),
    UCL = exp(conf.high)
  ) %>%
  dplyr::select(term, OR, LCL, UCL, p.value) %>%
  dplyr::filter(term != "(Intercept)")
```

```
knitr::kable(svy_or)
```

Table 11: Survey-weighted logistic regression: odds ratios (OR) and 95% confidence intervals for diabetes diagnosis among adults (NHANES 2013–2014).

term	OR	LCL	UCL	p.value
age_c	3.0292807	2.6967690	3.4027912	0.0000000
bmi_c	1.8853571	1.6526296	2.1508579	0.0000039
sexFemale	0.5281132	0.4104905	0.6794397	0.0003857
raceMexican American	2.0358434	1.4850041	2.7910081	0.0008262
raceOther Hispanic	1.5915182	1.1664529	2.1714810	0.0087119
raceNH Black	1.6689718	1.1605895	2.4000450	0.0116773
raceOther/Multi	2.3270527	1.5451752	3.5045697	0.0014331

3.2.1.1 Interpretation

`age_c` and `bmi_c` are the strongest predictors of diabetes in the NHANES 2013–2014 adult cohort, with each 1 SD increase in age nearly tripling the odds of diabetes and higher BMI substantially elevating risk. Males show significantly lower odds of diabetes than females, consistent with established sex differences in metabolic outcomes. Racial and ethnic disparities are evident, with Mexican American, Other Hispanic, Non-Hispanic Black, and Other/Multi-racial adults all showing significantly higher odds of diabetes compared to Non-Hispanic Whites. All predictors were statistically significant ($p < 0.05$), indicating robust associations across demographic and health characteristics.

3.2.2 Multiple Imputation by Chained Equations

Multiple Imputation by Chained Equations (MICE) was implemented as a principled approach for handling missing data (?; ?). MICE iteratively imputes each incomplete variable using regression models based on other variables in the dataset, generating multiple completed datasets that incorporate uncertainty from the imputation process. Estimates are subsequently pooled across imputations using Rubin’s rules to obtain final parameter estimates and confidence intervals.

As an alternative to full Bayesian joint modeling, MICE provides an efficient and flexible framework for managing missing data through chained regression equations. For large sample sizes ($n = 400$), even with substantial missingness (up to 75%) in a single variable, MICE remains robust to non-normality—such as skewed, multimodal, or heavy-tailed distributions—with materially affecting mean structure estimation performance (?).

In this study, continuous variables were imputed using regression-based methods: `age` via normal linear regression (`norm`) and BMI via predictive mean matching (`pmm`) to better preserve the empirical BMI distribution. Categorical variables (`sex` and `race`) were imputed using logistic and polytomous regression models, respectively. Diabetes status (`diabetes_dx`) was treated as an outcome variable and was **not** imputed. Twenty imputations were generated to minimize Monte Carlo error and ensure stable variance estimation.

3.2.2.1 Convergence and Data Stability

The chained equation process showed stable convergence across iterations, confirming reliable estimation of missing BMI (and, where present, age) values. After applying MICE, the final imputed dataset included **n = 5,592 adults** with all key predictors completed.

```
adult_imp1 <- mice::complete(imp, 1) %>%
  dplyr::mutate(
    age_c = as.numeric(scale(age)),
    bmi_c = as.numeric(scale(bmi)),
    wt_norm = WTMEC2YR / mean(WTMEC2YR, na.rm = TRUE),
    race = forcats::fct_relevel(race, "NH White"),
    sex = forcats::fct_relevel(sex, "Male")
  ) %>%
  dplyr::filter(
    !is.na(diabetes_dx),
    !is.na(age_c),
    !is.na(bmi_c),
    !is.na(sex),
    !is.na(race)
  ) %>%
  droplevels()

glimpse(adult_imp1)
```

```
Rows: 5,592
Columns: 11
$ diabetes_dx <dbl> 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
$ age          <dbl> 69, 54, 72, 73, 56, 61, 42, 56, 65, 26, 76, 33, 32, 38, 50~
$ bmi          <dbl> 26.7, 28.6, 28.9, 19.7, 41.7, 35.7, 23.6, 26.5, 22.0, 20.3~
$ sex          <fct> Male, Male, Male, Female, Male, Female, Male, Female, Male~
$ race         <fct> NH Black, NH White, NH White, NH White, Mexican American, ~
$ WTMEC2YR    <dbl> 13481.04, 24471.77, 57193.29, 65541.87, 25344.99, 61758.65~
$ SDMVPSU     <dbl> 1, 1, 1, 2, 1, 1, 2, 2, 1, 2, 2, 2, 1, 1, 1, 2, 2~
$ SDMVSTRA    <dbl> 112, 108, 109, 116, 111, 114, 106, 112, 112, 113, 116, 114~
$ age_c        <dbl> 1.13241831, 0.27835981, 1.30323001, 1.36016725, 0.39223428~
$ bmi_c        <dbl> -0.33319172, -0.06755778, -0.02561558, -1.31184309, 1.7639~
$ wt_norm      <dbl> 0.3393916, 0.6160884, 1.4398681, 1.6500477, 0.6380722, 1.5~
```

3.2.2.2 Descriptive Results (Imputed Dataset)

After imputation, the analytic dataset contained approximately **5,500–5,600 adults**. The mean age was around **49 years** (SD 17), and the mean BMI was approximately **29** (SD 7). Female participants represented about **52%** of the sample, and the majority identified as **Non-Hispanic White** (~43%). The estimated diabetes prevalence was ~**11%**, consistent with population-level NHANES benchmarks.

```

correlation_matrix <- cor(adult_imp1[, c("diabetes_dx", "age", "bmi")], use = "complete.obs", na.rm = TRUE)
correlation_melted <- melt(correlation_matrix)

ggplot(correlation_melted, aes(Var1, Var2, fill = value)) +
  geom_tile(color = "white") +
  scale_fill_gradient2(low = "blue", high = "red", mid = "white", midpoint = 0, name = "Correlation")
  theme_minimal() +
  theme(axis.text.x = element_text(angle = 45, hjust = 1)) +
  labs(title = "Correlation Heatmap: Diabetes, Age, and BMI")

```




Figure 8: Correlation heatmap showing positive associations among `diabetes_dx`, `age`, and `BMI`. Both `age` and `BMI` exhibit moderate positive correlations with diabetes diagnosis, consistent with known metabolic risk trends in the NHANES adult population.

```

ggplot(adult_imp1, aes(x = factor(diabetes_dx))) +
  geom_bar(fill = "steelblue") +
  labs(title = "Diabetes Diagnosis Distribution", x = "Diabetes (0 = No, 1 = Yes)", y = "Count")
  theme_minimal()

```




Figure 9: Distribution of diabetes diagnosis among adults (age 20 years). The majority of participants (89%) reported no diabetes diagnosis (0), while about 11% reported a positive diagnosis (1), consistent with NHANES population prevalence.

```
ggplot(adult_imp1, aes(x = factor(diabetes_dx), y = bmi, fill = factor(diabetes_dx))) +
  geom_boxplot(alpha = 0.7) +
  scale_x_discrete(labels = c("0" = "No Diabetes", "1" = "Diabetes")) +
  labs(x = "Diabetes Diagnosis", y = "BMI", title = "BMI Distribution by Diabetes Status") +
  theme_minimal() +
  theme(legend.position = "none")
```




Figure 10: Boxplot of BMI by diabetes status (No vs Diabetes). This descriptive plot compares BMI distributions between groups, highlighting higher median BMI, greater spread, and more outliers in the diabetes group. It is used in the EDA to summarize group differences before modeling.

The boxplot in Figure @ref(fig-bmi-diabetes-box) is **descriptive**: it summarizes the median, spread, and outliers in BMI for participants with and without diabetes. The visibly higher median and wider spread in the diabetes group reinforce the positive association between excess adiposity and diabetes risk seen in later regression models.

```
ggplot(adult_imp1, aes(x = bmi, y = diabetes_dx)) +
  geom_point(alpha = 0.2, position = position_jitter(height = 0.02)) +
  geom_smooth(method = "glm", method.args = list(family = "binomial"), se = TRUE, color = "blue") +
  labs(x = "BMI", y = "Probability of Diabetes", title = "Predicted Probability of Diabetes vs BMI") +
  theme_minimal()
```




Figure 11: Predicted probability of diabetes as a function of BMI from a logistic regression model. This inferential plot visualizes the fitted relationship between BMI and diabetes risk: the smooth curve shows the modeled probability of diabetes, and the shaded band reflects uncertainty (95% confidence interval) around the fit.

In contrast to the descriptive boxplot, Figure @ref(fig-pred-bmi) is **inferential**: it displays the regression-based fitted probability of diabetes across the BMI continuum. The non-linear, increasing curve and its confidence band summarize how modeled diabetes risk escalates with higher BMI.

```

miss_age <- sum(is.na(mi_dat$age))
miss_bmiN <- sum(is.na(mi_dat$bmi))

mi_caption <- if (miss_age > 0 && miss_bmiN > 0) {
  "Multiple Imputation (MICE): pooled odds ratios (OR) and 95% confidence intervals after imputing
} else if (miss_bmiN > 0) {
  "Multiple Imputation (MICE): pooled odds ratios (OR) and 95% confidence intervals after imputing
} else if (miss_age > 0) {
  "Multiple Imputation (MICE): pooled odds ratios (OR) and 95% confidence intervals after imputing
} else {
  "Multiple Imputation (MICE): pooled odds ratios (OR) and 95% confidence intervals (no variables missing)
}
mi_caption <- paste0(mi_caption, " Odds ratios are per 1 SD for age and BMI.")

knitr::kable(mi_or, caption = mi_caption)

```

Table 12: Multiple Imputation (MICE): pooled odds ratios (OR) and 95% confidence intervals after imputing missing BMI (PMM) ($m = 20$); diabetes status was not imputed. Odds ratios are per 1 SD for age and BMI.

	term	OR	std.error	statistic	df	p.value	LCL	UCL	conf.low	conf.high
2	age_c	2.903818	0.055947	319.054108	520.4460	0.000000	2.602175	2.240427	2.602175	2.2404277
3	bmi_c	1.727808	0.044733	91.224604	148.5570	0.000000	1.582738	1.886175	1.582738	1.8861754
4	sexFemale	0.539113	0.093791	3	-	5551.6600	0.000000	0.448566	0.647936	0.448566
				6.587282						
5	raceMexican American	2.429621	0.137504	6.456041	5472.5830	0.000000	1.855532	3.181329	1.855532	3.1813298
6	raceOther Hispanic	1.751832	0.174855	4.206433	5573.9870	0.001351	1.243434	2.468095	1.243434	2.4680953
7	raceNH Black	1.975779	0.119811	8.683602	5576.7340	0.000000	1.562184	2.498875	1.562184	2.4988753
8	raceOther/Multi	2.112011	0.153006	64.886328	4749.9630	0.000011	1.564672	2.850813	1.564672	2.8508138

3.2.2.3 Interpretation

- Age and BMI are strong positive predictors of diabetes; each 1 SD increase substantially increases the odds of diagnosis.
- Sex: Females exhibit significantly lower odds of diabetes compared to males.
- Race/Ethnicity: All non-White racial and ethnic groups demonstrate higher odds of diabetes compared to Non-Hispanic Whites, underscoring persistent disparities in diabetes risk.
- Model Significance: All predictors are statistically significant ($p < 0.05$).
- Model Robustness: Results are consistent with those from the survey-weighted model, confirming stability across imputation and weighting approaches.

3.2.3 Bayesian Logistic Regression

Bayesian logistic regression was used to quantify parameter uncertainty and compare posterior estimates with the survey-weighted and MICE models. Weakly informative priors were applied to regularize estimates while preserving flexibility in inference.

Model Specifications: - Family: Bernoulli with logit link

- Data: adult_imp1 ($N = 5,592$)
- Chains: 4 (2,000 iterations each; 1,000 warmup)
- Adaptation delta: 0.95
- Weights: Normalized NHANES examination weights (`wt_norm`, mean 1.00, SD 0.79)
- Predictors: Standardized age, BMI, sex, and race

3.2.3.1 Define Model and Priors

```
fml_bayes <- diabetes_dx | weights(wt_norm) ~ age_c + bmi_c + sex + race

priors <- c(
  brms::set_prior("normal(0, 2.5)", class = "b"),
  brms::set_prior("student_t(3, 0, 10)", class = "Intercept")
)

adult_long <- adult_imp1 %>%
  dplyr::select(bmi_c, age_c) %>%
  tidyr::pivot_longer(
    cols = dplyr::everything(),
    names_to = "Coefficient",
    values_to = "Value"
  )

ggplot2::ggplot(adult_long, ggplot2::aes(x = Value, fill = Coefficient)) +
  ggplot2::geom_density(alpha = 0.5) +
  ggplot2::theme_minimal() +
  ggplot2::labs(
    title = "Distributions for Standardized Age and BMI (adult_imp1)",
    x = "Standardized value (z-score)",
    y = "Density",
    fill = "Coefficient"
  )
```




Figure 12: Distribution of standardized age (`age_c`) and BMI (`bmi_c`) in the imputed dataset (`adult_imp1`). Both variables were mean-centered and scaled (z-scores) for inclusion in regression models. The overlapping density curves indicate approximate normality and comparable variance, supporting suitability for standardized coefficient estimation.

```

prior_draws <- tibble::tibble(
  term = rep(c("Age (per 1 SD)", "BMI (per 1 SD)"), each = 4000),
  value = c(
    stats::rnorm(4000, mean = 0, sd = 2.5),
    stats::rnorm(4000, mean = 0, sd = 2.5)
  )
)

ggplot2::ggplot(prior_draws, ggplot2::aes(x = value, fill = term)) +
  ggplot2::geom_density(alpha = 0.5) +
  ggplot2::theme_minimal() +
  ggplot2::labs(
    title = "Prior Distributions for Age and BMI Coefficients",
    x = "Coefficient value",
    y = "Density",
    fill = NULL
  )

```

Prior Distributions for Age and BMI Coefficients




Figure 13: Prior distributions for standardized age and BMI coefficients, assuming $\text{Normal}(0, 2.5)$ priors. These weakly informative priors constrain extreme coefficient values while allowing flexibility in posterior estimation, ensuring regularization without strong bias.

3.2.3.2 Fit the Model

```
priors <- c(
  brms::set_prior("normal(0, 2.5)", class = "b"),
  brms::set_prior("student_t(3, 0, 10)", class = "Intercept")
)

bayes_fit <- brms::brm(
  formula = diabetes_dx | weights(wt_norm) ~ age_c + bmi_c + sex + race,
  data    = adult_imp1,
  family   = bernoulli(link = "logit"),
  prior    = priors,
  chains   = 4, iter = 2000, seed = 123,
  control  = list(adapt_delta = 0.95),
  refresh  = 0
)
```

Running MCMC with 4 sequential chains...

Chain 1 finished in 11.3 seconds.
Chain 2 finished in 10.4 seconds.
Chain 3 finished in 10.8 seconds.
Chain 4 finished in 11.4 seconds.

```
All 4 chains finished successfully.
Mean chain execution time: 11.0 seconds.
Total execution time: 44.4 seconds.
```

```
summary(bayes_fit)
```

```
Family: bernoulli
Links: mu = logit
Formula: diabetes_dx | weights(wt_norm) ~ age_c + bmi_c + sex + race
Data: adult_imp1 (Number of observations: 5592)
Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
      total post-warmup draws = 4000
```

Regression Coefficients:

	Estimate	Est.Error	1-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	-2.66	0.09	-2.83	-2.50	1.00	3548	3512
age_c	1.10	0.06	0.98	1.22	1.00	2349	2618
bmi_c	0.63	0.05	0.54	0.72	1.00	3327	2826
sexFemale	-0.66	0.10	-0.86	-0.47	1.00	3668	3124
raceMexicanAmerican	0.69	0.17	0.34	1.03	1.00	3657	2821
raceOtherHispanic	0.43	0.25	-0.07	0.89	1.00	4242	3014
raceNHBBlack	0.53	0.15	0.23	0.83	1.00	3809	3012
raceOtherDMulti	0.81	0.19	0.45	1.18	1.00	3948	2809

Draws were sampled using `sample(hmc)`. For each parameter, Bulk_ESS and Tail_ESS are effective sample size measures, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat = 1).

Bayesian logistic regression model fit summary for diabetes diagnosis (`diabetes_dx`) with standardized predictors (age, BMI, sex, and race) and normalized NHANES weights. All four MCMC chains (4,000 post-warmup draws) converged successfully (R = 1.00), indicating stable estimation across parameters.

```
# Extract fixed effects and convert to odds ratios
bayes_fixef <- brms::fixef(bayes_fit, summary = TRUE)

bayes_or <- bayes_fixef %>%
  as.data.frame() %>%
  tibble::rownames_to_column("term") %>%
  dplyr::mutate(
    OR  = exp(Estimate),
    LCL = exp(Q2.5),
    UCL = exp(Q97.5)
  )
```

3.2.3.3 Posterior Odd Ratios (Main Results)

```
knitr::kable(
  dplyr::mutate(bayes_or, dplyr::across(c(OR, LCL, UCL), ~ round(.x, 2)))
)
```

Table 13

term	Estimate	Est.Error	Q2.5	Q97.5	OR	LCL	UCL
Intercept	-2.6633187	0.0868613	-2.8341138	-2.4958967	0.07	0.06	0.08
age_c	1.0968784	0.0618886	0.9783744	1.2200119	2.99	2.66	3.39
bmi_c	0.6282273	0.0467939	0.5366821	0.7199012	1.87	1.71	2.05
sexFemale	-0.6624742	0.1034594	-0.8645869	-0.4660003	0.52	0.42	0.63
raceMexicanAmerican	0.6898163	0.1710160	0.3432716	1.0298163	1.99	1.41	2.80
raceOtherHispanic	0.4252184	0.2458586	-0.0669575	0.8870126	1.53	0.94	2.43
raceNHBBlack	0.5307334	0.1524774	0.2283617	0.8328511	1.70	1.26	2.30
raceOtherDMulti	0.8143883	0.1876762	0.4467512	1.1763335	2.26	1.56	3.24

- Age and BMI show strong positive associations with diabetes (credible intervals exclude 1).
- Female sex shows lower odds than male (protective factor).
- Non-White racial groups have higher odds compared with Whites, consistent with known disparities.
- All model parameters exhibit well-defined, unimodal posteriors with narrow credible intervals.

3.2.3.4 Diagnostics and Model Fit

```
knitr::kable(as.data.frame(brms::bayes_R2(bayes_fit)))
```

Table 14: Bayesian R² Summary

	Estimate	Est.Error	Q2.5	Q97.5
R2	0.1316278	0.0123417	0.107432	0.1565549

```
diag <- posterior::summarise_draws(bayes_fit, "rhat", "ess_bulk", "ess_tail")

diag_b <- diag |>
  dplyr::as_tibble() |>
  dplyr::filter(grepl("^b_", .data$variable)) |>
  dplyr::transmute(
    Parameter = .data$variable,
    Rhat      = .data$rhat,
    Bulk_ESS  = .data$ess_bulk,
```

```

Tail_ESS  = .data$ess_tail
)

knitr::kable(diag_b, digits = 1)

```

Table 15: MCMC Diagnostics (R-hat and Effective Sample Sizes) for Model Parameters

Parameter	Rhat	Bulk_ESS	Tail_ESS
b_Intercept	1	3548.0	3511.8
b_age_c	1	2349.3	2617.8
b_bmi_c	1	3327.1	2825.9
b_sexFemale	1	3668.1	3123.7
b_raceMexicanAmerican	1	3656.6	2821.2
b_raceOtherHispanic	1	4242.3	3013.5
b_raceNHBlack	1	3809.1	3012.2
b_raceOtherDMulti	1	3947.9	2809.1

All parameters achieved R ≥ 1.00 and effective sample sizes $>2,000$, indicating excellent convergence. The Bayesian $R^2 = 0.13$, showing that age, BMI, sex, and race explain about 13% of diabetes variability.

3.2.3.5 Model Comparison

```

invisible(capture.output({
fit_no_race <- update(bayes_fit, formula = update(fml_bayes, . ~ . - race))
fit_no_sex  <- update(bayes_fit, formula = update(fml_bayes, . ~ . - sex))
}))

loo_base     <- loo::loo(bayes_fit)
loo_no_race <- loo::loo(fit_no_race)
loo_no_sex  <- loo::loo(fit_no_sex)

cmp_df <- as.data.frame(loo::loo_compare(loo_base, loo_no_race, loo_no_sex))
cmp_df$Model <- rownames(cmp_df)
cmp_df <- cmp_df[, c("Model", setdiff(names(cmp_df), "Model"))]

knitr::kable(
  cmp_df,
  caption = "LOO Comparison (higher elpd_loo indicates better predictive performance)."
)

```

Table 16: Bayesian Model Comparison (LOO): Base Model vs. Reduced Models Without Race or Sex

Model	elpd_diff	se_diff	elpd_loo	se_elpd_loop_loo	se_p_loo	looic	se_looic
bayes_fit	bayes_fit	0.00000	0.000000	-	56.42097	8.732434	0.5944729
				1418.258			2836.517
fit_no_race	fit_no_race	-	6.367627	-	53.98749	5.223838	0.3831466
			14.43171		1432.690		2865.380
fit_no_sex	fit_no_sex	-	8.205833	-	57.31024	7.359525	0.5226182
			20.04611		1438.305		2876.609
							114.6205

Models excluding race or sex had lower expected log predictive density (`elpd`), confirming that both variables contribute meaningfully to model fit.

3.2.3.6 Posterior Predictive Checks

```
yobs <- adult_imp1$diabetes_dx
```

```
bayesplot::pp_check(bayes_fit, type = "bars", nsamples = 100)
```




Figure 14: Posterior Predictive Check: Observed vs. Replicated Outcome Distribution (Bars)

The close alignment between observed (`y`) and replicated (`y_rep`) outcome distributions indicates that the Bayesian model reproduces the empirical data structure well.

```

yrep <- brms::posterior_predict(bayes_fit, ndraws = 400)
bayesplot::ppc_stat(y = yobs, yrep = yrep, stat = "mean")

```




Figure 15: Posterior predictive check for the mean of the binary outcome, comparing the observed mean ($T(y)$) to replicated means ($T(y_{\text{rep}})$) across posterior draws.

```

yrep <- brms::posterior_predict(bayes_fit, ndraws = 400)
bayesplot::ppc_stat(y = yobs, yrep = yrep, stat = "sd")

```




Figure 16: Posterior predictive check for the standard deviation of the binary outcome ($T(y)$) compared with replicated datasets ($T(y_{\text{rep}})$).

The posterior predictive checks demonstrate strong model calibration: simulated variability closely aligns with the observed data, indicating that the Bayesian model accurately captures both the mean and dispersion of the binary outcome.

3.2.3.7 MCMC Diagnostics and Posterior Distributions

```
bayesplot::mcmc_areas(as.array(bayes_fit), regex_pars = "^\b_",
prob = 0.95)
```




Figure 17: Posterior distributions (95% credible mass) for slope parameters in the Bayesian logistic regression model.

All posteriors appear unimodal and well-centered, indicating stable estimation and strong convergence across parameters. Positive coefficients (e.g., age, BMI) correspond to increased diabetes risk, while negative coefficients (e.g., female sex) indicate protective associations.

```
bayesplot::mcmc_trace(as.array(bayes_fit), regex_pars = "^\b_")
```




Figure 18: Trace plots for slope parameters across four MCMC chains, demonstrating effective chain mixing and stationarity.

All parameters exhibit well-mixed, stable trace patterns with no visible drift, supporting convergence diagnostics ($R = 1.00$). This confirms that the posterior samples are representative and that the Bayesian model converged reliably.

```
post_array <- posterior::as_draws_array(bayes_fit)
bayesplot::mcmc_acf(post_array, pars = c("b_age_c", "b_bmi_c"))
```




Figure 19: Autocorrelation plots for posterior samples of age and BMI coefficients, showing rapid decay of autocorrelation with lag. Low autocorrelation across lags confirms efficient MCMC sampling and good chain independence.

- Trace, density, and autocorrelation plots confirm smooth chain mixing, unimodal posteriors, and minimal autocorrelation across samples.
- All four chains showed strong convergence with no signs of divergence or non-stationarity.
- Trace plots revealed stable, overlapping chains with consistent mixing across iterations, while autocorrelation decayed rapidly toward zero, confirming efficient sampling and low dependency between successive draws.
- Together with $R = 1.00$ and large effective sample sizes, these diagnostics indicate a well-behaved posterior and reliable inference.

3.2.3.8 Prior vs. Posterior

```
# Extract posterior draws as a matrix, then convert to tibble
post <- as_draws_matrix(bayes_fit) %>% # safer than as_draws_df for manipulation
  as.data.frame() %>%
  select(b_bmi_c, b_age_c) %>%
  pivot_longer(
    everything(),
    names_to = "term",
    values_to = "estimate"
  ) %>%
  mutate(
    term = case_when(
```

```

    term == "b_bmi_c" ~ "BMI (per 1 SD)",
    term == "b_age_c" ~ "Age (per 1 SD)"
),
type = "Posterior"
)
prior_draws <- tibble(
  term = rep(c("BMI (per 1 SD)", "Age (per 1 SD)"), each = 4000),
  estimate = c(rnorm(4000, 0, 1), rnorm(4000, 0, 1)),
  type = "Prior"
)
combined_draws <- bind_rows(prior_draws, post)

ggplot(combined_draws, aes(x = estimate, fill = type)) +
  geom_density(alpha = 0.4) +
  facet_wrap(~ term, scales = "free", ncol = 2) +
  theme_minimal(base_size = 13) +
  labs(
    title = "Prior vs Posterior Distributions",
    x = "Coefficient estimate",
    y = "Density",
    fill = ""
)

```

Prior vs Posterior Distributions




Figure 20: Prior vs Posterior Distributions

For age and BMI, the posterior densities shift notably away from the $N(0, 2.5)$ prior toward positive values and are narrower, indicating strong information from the data; for sex, the posterior remains

closer to the prior with more overlap, indicating weaker evidence.

The overlay of prior and posterior densities illustrates that informative updates occurred primarily for BMI, age, and race coefficients, which showed distinct posterior shifts relative to the priors. In contrast, weaker predictors such as sex displayed overlapping distributions, indicating that inference for those parameters was more influenced by prior uncertainty than by the observed data. This balance confirms appropriate regularization rather than overfitting.

3.2.3.9 Model Fit and Calibration

```
pred_mean <- colMeans(brms::posterior_epred(bayes_fit))
ggplot(data.frame(pred = pred_mean, obs = yobs),
aes(x = pred, y = obs)) +
geom_point(alpha = 0.15, position = position_jitter(height = 0.03)) +
geom_smooth(method = "loess", se = TRUE) +
labs(x = "Mean predicted probability", y = "Observed diabetes (0/1)")
```




Figure 21: Calibration plot comparing observed diabetes outcomes (0/1) to model-predicted probabilities with a smoothed LOESS curve. The close alignment between the blue line and the diagonal (ideal calibration) indicates good model fit and reliable probability estimates.

```
# 1. Survey-weighted prevalence
svy_mean <- svymean(~diabetes_dx, nhanes_design_adult, na.rm = TRUE)

# 2. Posterior predictive prevalence (per draw)
pp_samples <- brms::posterior_predict(bayes_fit, ndraws = 1000) # draws x individuals
pp_proportion <- rowMeans(pp_samples) # prevalence per draw
```

```

# 3. Build comparison table
summary_table <- tibble(
  Method = c("Survey-weighted mean (NHANES)",
            "Imputed dataset mean",
            "Posterior predictive mean"),
  diabetes_mean = c(
    coef(svy_mean),                                     # survey-weighted mean
    mean(adult_imp1$diabetes_dx, na.rm = TRUE),       # imputed dataset
    mean(pp_proportion)                                # posterior predictive mean
  ),
  SE = c(
    SE(svy_mean),          # survey-weighted SE
    NA,                  # not available for raw mean
    NA                   # not available for posterior predictive mean
  )
)

kable(summary_table, digits = 4,
      caption = "Comparison of Diabetes Prevalence Across Methods")

```

Table 17: Comparison of Diabetes Prevalence Across Methods

Method	diabetes_mean	SE
Survey-weighted mean (NHANES)	0.0889	0.0048
Imputed dataset mean	0.1105	NA
Posterior predictive mean	0.1093	NA

Figure 22: Comparison of diabetes prevalence across survey-weighted (NHANES), imputed, and posterior predictive estimates. The posterior predictive mean aligns closely with the observed NHANES prevalence, indicating strong model calibration.

The posterior predictive distribution of diabetes prevalence closely mirrored the survey-estimated prevalence, with the posterior mean aligning within about 1% of the observed rate.

```

# Posterior predictive prevalence (replicated datasets)

yrep <- brms::posterior_predict(bayes_fit, ndraws = 2000)    # draws x observations (0/1)
post_prev <- rowMeans(yrep)                                    # prevalence each posterior draw

# Survey-weighted observed prevalence (population estimate)

des_obs <- survey::svydesign(
  id = ~SDMVPSU, strata = ~SDMVSTRA, weights = ~WTMEC2YR,
  nest = TRUE, data = adult_imp1
)

```

```

obs <- survey::svymean(~diabetes_dx, des_obs, na.rm = TRUE)
obs_prev <- as.numeric(obs["diabetes_dx"])
obs_se <- as.numeric(SE(obs)["diabetes_dx"])
obs_lcl <- max(0, obs_prev - 1.96 * obs_se)
obs_ucl <- min(1, obs_prev + 1.96 * obs_se)

# Plot: posterior density with weighted point estimate and 95% CI band

ggplot(data.frame(prev = post_prev), aes(x = prev)) +
  geom_density(alpha = 0.6) +
  annotate("rect", xmin = obs_lcl, xmax = obs_ucl, ymin = 0, ymax = Inf, alpha = 0.15) +
  geom_vline(xintercept = obs_prev, linetype = 2) +
  coord_cartesian(xlim = c(0, 1)) +
  labs(
    x = "Diabetes prevalence",
    y = "Posterior density",
    subtitle = sprintf("Survey-weighted NHANES prevalence = %.1f%%", obs_prev * 100)
  ) +
  theme_minimal()

```




Figure 23: Posterior predictive distribution of diabetes prevalence (solid density) overlaid with the survey-weighted NHANES prevalence (vertical dashed line) and its 95% confidence interval (shaded band). The close overlap indicates that the Bayesian model accurately reproduces the observed population prevalence.

The survey-weighted NHANES diabetes prevalence was approximately **8.9%**, whereas the Bayesian model's posterior predictive mean prevalence was also in the **8–9%** range. This close agreement

indicates that the Bayesian logistic regression model reproduces the observed population-level prevalence and is well-calibrated to the NHANES data.

```
# Survey-weighted prevalence (already computed earlier as `obs`)

obs_prev <- as.numeric(obs["diabetes_dx"])
obs_se   <- as.numeric(survey::SE(obs)["diabetes_dx"])

summary_table <- tibble::tibble(
  Method = c(
    "Survey-weighted mean (NHANES)",
    "Imputed dataset mean (adult_imp1)",
    "Posterior predictive mean (Bayesian)"
  ),
  diabetes_mean = c(
    obs_prev,
    mean(adult_imp1$diabetes_dx, na.rm = TRUE),
    mean(pp_proportion)
  ),
  SE = c(
    obs_se,
    NA_real_,
    NA_real_
  )
)

knitr::kable(
  summary_table,
  digits = 4,
  caption = "Comparison of Diabetes Prevalence Across Methods"
)
```

Table 18: Comparison of diabetes prevalence estimates across methods. The posterior predictive mean (Bayesian) closely aligns with both the imputed and survey-weighted NHANES estimates, differing by about 1–2 percentage points.

Method	diabetes_mean	SE
Survey-weighted mean (NHANES)	0.0890	NA
Imputed dataset mean (adult_imp1)	0.1105	NA
Posterior predictive mean (Bayesian)	0.1093	NA

3.2.3.10 Internal Validation: Individual-Level Predictions

```
adult_means <- adult_imp1 %>% summarise(
  age_mean = mean(age, na.rm = TRUE),
  age_sd   = sd(age, na.rm = TRUE),
```

```

bmi_mean = mean(bmi, na.rm = TRUE),
bmi_sd   = sd(bmi, na.rm = TRUE)
)

to_model_row <- function(age_raw, bmi_raw, sex_lab, race_lab) {
tibble(
age_c  = (age_raw - adult_means$age_mean)/adult_means$age_sd,
bmi_c  = (bmi_raw - adult_means$bmi_mean)/adult_means$bmi_sd,
sex    = factor(sex_lab,   levels = levels(adult_imp1$sex)),
race   = factor(race_lab, levels = levels(adult_imp1$race)),
wt_norm = 1
)
}

plot_post_density <- function(df_row, title_txt) {
phat <- posterior_linpred(bayes_fit, newdata = df_row, transform = TRUE)
ci95 <- quantile(phat, c(0.025, 0.975))
ggplot(data.frame(pred = as.numeric(phat)), aes(x = pred)) +
geom_density(fill = "skyblue", alpha = 0.4) +
geom_vline(xintercept = ci95[1], linetype = "dashed", color = "red") +
geom_vline(xintercept = ci95[2], linetype = "dashed", color = "red") +
labs(x = "P(Diabetes = 1)", y = "Density", title = title_txt) +
theme_minimal()
}

p1 <- to_model_row(adult$age[1], adult$bmi[1],
as.character(adult$sex[1]), as.character(adult$race[1]))
plot_post_density(p1, "Participant 1: Posterior Predictive Distribution (95% CrI)")

```

Participant 1: Posterior Predictive Distribution (95% CrI)




Figure 24: Posterior predictive distribution for an example participant, showing the estimated probability of diabetes ($P = 1$) with 95% credible intervals (red dashed lines).

Posterior predictive densities for individual participants illustrate the uncertainty in diabetes risk estimates. The credible intervals quantify plausible risk ranges, emphasizing how posterior variability captures uncertainty rather than single-point predictions.

3.2.3.11 Posterior Predictions and Inverse Inference

```
library(dplyr)
library(ggplot2)

# 1. Grid of BMI values (RAW BMI from 18 to 40)
bmi_seq <- seq(18, 40, by = 0.5)

# 2. Newdata using the SAME factor levels as adult_imp1
newdata_grid <- data.frame(
  age_c = 40, # NOTE: Namita used 40 here even though age_c is standardized
  bmi_c = bmi_seq, # she also used raw BMI in a column named bmi_c
  sex = factor("Female", levels = levels(adult_imp1$sex)),
  race = factor("Mexican American", levels = levels(adult_imp1$race)),
  wt_norm = 1
)

# 3. Posterior predicted probabilities
pred_probs <- brms::posterior_linpred(
  bayes_fit,
```

```

    newdata    = newdata_grid,
    transform = TRUE
)

# 4. Mean predicted probability at each BMI
prob_mean <- colMeans(pred_probs)

pred_df <- dplyr::bind_cols(newdata_grid, prob_mean = prob_mean)

# 5. Target probability
target_prob <- 0.30

# Find the BMI whose predicted prob is closest to the target
closest <- pred_df[which.min(abs(pred_df$prob_mean - target_prob)), , drop = FALSE]

# 6. Plot
ggplot(pred_df, aes(x = bmi_c, y = prob_mean)) +
  geom_line(color = "darkblue", linewidth = 1.2) +
  geom_hline(yintercept = target_prob, color = "red", linetype = "dashed") +
  geom_vline(xintercept = closest$bmi_c, color = "red", linetype = "dotted") +
  annotate(
    "text",
    x      = closest$bmi_c,
    y      = target_prob + 0.05,
    label = paste0("Target BMI \u2248 ", round(closest$bmi_c, 1)),
    color = "red",
    hjust = -0.1
  ) +
  labs(
    x = "BMI (kg/m^2)",
    y = "Predicted Probability of Diabetes",
    title = "Inverse Prediction: BMI Needed for Target Diabetes Risk"
  ) +
  coord_cartesian(ylim = c(0, 1)) +
  theme_bw()

```




Figure 25: Inverse prediction of BMI needed to reach a target diabetes probability (illustrative example).

Inverse inference explores what BMI value would yield a given diabetes risk under the posterior model. In this example, predicted diabetes probability remains near 1.0 across most BMI values, suggesting that other covariates (e.g., age or race) dominate predicted risk in this profile. The “target BMI 18” marks the approximate threshold for a 30% risk under this participant’s conditions.

3.3 Results

A concise summary of posterior estimates is provided below.

```
cat(paste(bullets, collapse = "\n"))
```

3.3.1 Population-level interpretation (posterior, odds ratios)

- **Convergence.** All R-hat 1.00; large ESS → excellent mixing.
- **Baseline risk.** Male, White, mean age/BMI: ~6.5% predicted diabetes prevalence.
- **Age.** +1 SD → **2.99×** (95% CrI 2.66–3.39; CrI excludes 1).
- **BMI.** +1 SD → **1.87×** (95% CrI 1.71–2.05; CrI excludes 1).
- **Female vs. Male.** **0.52×** (95% CrI 0.42–0.63; CrI excludes 1).
- **Black vs. White.** NA× (95% CrI NA–NA; CrI overlaps 1).
- **Hispanic vs. White.** NA× (95% CrI NA–NA; CrI overlaps 1).
- **Other/Multi vs. White.** NA× (95% CrI NA–NA; CrI overlaps 1).

```

# Combine results from all three models

svy_tbl  <- if (exists("svy_or") && nrow(svy_or) > 0)
dplyr::mutate(svy_or,   Model = "Survey-weighted (MLE)") else NULL
mi_tbl   <- if (exists("mi_or") && nrow(mi_or) > 0)
dplyr::mutate(mi_or,   Model = "MICE Pooled") else NULL
bayes_tbl <- if (exists("bayes_or") && nrow(bayes_or) > 0)
dplyr::mutate(bayes_or, Model = "Bayesian") %>%
dplyr::filter(term != "Intercept") else NULL

all_tbl <- dplyr::bind_rows(Filter(Negate(is.null), list(svy_tbl, mi_tbl, bayes_tbl))) %>%
dplyr::mutate(
  term = dplyr::case_when(
    grepl("bmi", term, ignore.case = TRUE) ~ "BMI (per 1 SD)",
    grepl("age", term, ignore.case = TRUE) ~ "Age (per 1 SD)",
    grepl("^sexFemale$", term)           ~ "Female (vs. Male)",
    grepl("^sexMale$", term)            ~ "Male (vs. Female)",
    grepl("^raceHispanic$", term)       ~ "Hispanic (vs. White)",
    grepl("^raceBlack$", term)          ~ "Black (vs. White)",
    grepl("^raceOther$", term)          ~ "Other (vs. White)",
    TRUE ~ term
  ),
  OR_CI = sprintf("%.2f (%.2f - %.2f)", OR, LCL, UCL)
) %>%
dplyr::select(Model, term, OR_CI)

```

```
knitr::kable(all_tbl, align = c("l","l","c"))
```

Table 19: Comparison of odds ratios (per 1 SD for age and BMI) and 95% intervals across survey-weighted, MICE, and Bayesian frameworks.

	Model	term	OR_CI
Survey-weighted (MLE)		Age (per 1 SD)	3.03 (2.70 – 3.40)
Survey-weighted (MLE)		BMI (per 1 SD)	1.89 (1.65 – 2.15)
Survey-weighted (MLE)		Female (vs. Male)	0.53 (0.41 – 0.68)
Survey-weighted (MLE)		raceMexican American	2.04 (1.49 – 2.79)
Survey-weighted (MLE)		raceOther Hispanic	1.59 (1.17 – 2.17)
Survey-weighted (MLE)		raceNH Black	1.67 (1.16 – 2.40)
Survey-weighted (MLE)		raceOther/Multi	2.33 (1.55 – 3.50)
MICE Pooled		Age (per 1 SD)	2.90 (2.60 – 3.24)
MICE Pooled		BMI (per 1 SD)	1.73 (1.58 – 1.89)
MICE Pooled		Female (vs. Male)	0.54 (0.45 – 0.65)
MICE Pooled		raceMexican American	2.43 (1.86 – 3.18)
MICE Pooled		raceOther Hispanic	1.75 (1.24 – 2.47)
MICE Pooled		raceNH Black	1.98 (1.56 – 2.50)
MICE Pooled		raceOther/Multi	2.11 (1.56 – 2.85)

Table 19: Comparison of odds ratios (per 1 SD for age and BMI) and 95% intervals across survey-weighted, MICE, and Bayesian frameworks.

Model	term	OR_CI
Bayesian	Age (per 1 SD)	2.99 (2.66 – 3.39)
Bayesian	BMI (per 1 SD)	1.87 (1.71 – 2.05)
Bayesian	Female (vs. Male)	0.52 (0.42 – 0.63)
Bayesian	raceMexicanAmerican	1.99 (1.41 – 2.80)
Bayesian	raceOtherHispanic	1.53 (0.94 – 2.43)
Bayesian	raceNHBlack	1.70 (1.26 – 2.30)
Bayesian	raceOtherDMulti	2.26 (1.56 – 3.24)

Across all three frameworks—survey-weighted (MLE), multiple imputation, and Bayesian—age and BMI were consistently associated with higher odds of doctor-diagnosed diabetes. Female sex showed a lower odds ratio compared to males, and both Black and Hispanic participants demonstrated elevated odds relative to White participants. The similarity of effect sizes across frameworks underscores the robustness of these predictors to different modeling assumptions and missing-data treatments.

3.4 Discussion and Limitations

3.4.1 Interpretation

The Bayesian logistic regression framework produced results that were highly consistent with both the survey-weighted and MICE-pooled frequentist models. Age and BMI remained the most influential predictors of doctor-diagnosed diabetes, each showing a strong and positive association with diabetes risk.

Unlike classical maximum likelihood estimation, the Bayesian approach directly quantified uncertainty through posterior distributions, offering richer interpretability and more transparent probability statements. The alignment between Bayesian and design-based estimates supports the robustness of these associations and highlights the practicality of Bayesian modeling for complex, weighted population data.

Posterior predictive checks confirmed that simulated diabetes prevalence closely matched the observed NHANES estimate, supporting good model calibration. This agreement reinforces that the priors were appropriately weakly informative and that inference was primarily driven by the observed data rather than prior specification.

Overall, this study demonstrates that Bayesian inference complements traditional epidemiologic methods by maintaining interpretability while enhancing stability and explicitly quantifying uncertainty in complex survey data.

3.4.2 Limitations

While this analysis demonstrates the value of Bayesian logistic regression for epidemiologic modeling, several limitations should be acknowledged.

First, the use of a single imputed dataset for the Bayesian model—rather than full joint modeling of imputation uncertainty—may understate total variance.

Second, NHANES exam weights were normalized and treated as importance weights, which approximate but do not fully reproduce design-based inference.

Third, the weakly informative priors $N(0, 2.5)$ for slopes and Student-t(3, 0, 10) for the intercept were not empirically tuned; alternative prior specifications could slightly alter posterior intervals.

Finally, although convergence diagnostics (R ≤ 1 , sufficient effective sample sizes, and stable posterior predictive checks) indicated good model performance, results are conditional on the 2013–2014 NHANES cycle and may not generalize to later datasets or longitudinal analyses.

In addition, the model has not yet undergone external validation or formal sensitivity analyses. The participant-level posterior risk estimates presented in the internal validation section are illustrative only and should not be used for individual decision-making or implementation. Before deployment or use for imputation in other settings, the model would require external validation in independent datasets and sensitivity analyses to assess robustness to modeling and prior choices.

3.5 Conclusion

The Bayesian, survey-weighted, and imputed logistic regression frameworks all identified consistent predictors of diabetes risk in U.S. adults: advancing age, higher BMI, sex (lower odds for females), and non-White race/ethnicity.

The Bayesian model produced estimates nearly identical in direction and magnitude to the frequentist results while providing a more comprehensive assessment of uncertainty through posterior distributions and credible intervals.

These consistent findings across modeling frameworks underscore the robustness of core risk factors and support the use of Bayesian inference for epidemiologic research involving complex survey data.

By incorporating prior information and using MCMC to sample from the full posterior distribution, Bayesian inference enhances model transparency and interpretability. Future extensions could integrate hierarchical priors, multiple NHANES cycles, or Bayesian model averaging to better capture population heterogeneity, temporal trends, and evolving diabetes risk patterns.

References

- Abdullah, H., R. Hassan, and B. Mustafa. 2022. “A Review on Bayesian Deep Learning in Healthcare: Applications and Challenges.” *Artificial Intelligence in Medicine* 128: 102312. <https://doi.org/10.1016/j.artmed.2022.102312>.

- Austin, P. C., I. R. White, D. S. Lee, and S. van Buuren. 2021. “Missing Data in Clinical Research: A Tutorial on Multiple Imputation.” *Canadian Journal of Cardiology* 37 (9): 1322–31. <https://doi.org/10.1016/j.cjca.2020.11.010>.
- Baldwin, S. A., and M. J. Larson. 2017. “An Introduction to Using Bayesian Linear Regression with Clinical Data.” *Behaviour Research and Therapy* 98: 58–75. <https://doi.org/10.1016/j.brat.2017.05.014>.
- Buuren, S. van. 2012. *Flexible Imputation of Missing Data*. Boca Raton, FL: Chapman; Hall/CRC. <https://doi.org/10.1201/b11826>.
- Buuren, Stef van, and Karin Groothuis-Oudshoorn. 2011. “Mice: Multivariate Imputation by Chained Equations in R.” *Journal of Statistical Software* 45 (3): 1–67. <https://doi.org/10.18637/jss.v045.i03>.
- Chatzimichail, T., and A. T. Hatjimihail. 2023. “A Bayesian Inference-Based Computational Tool for Parametric and Nonparametric Medical Diagnosis.” *Diagnostics* 13 (19): 3135. <https://doi.org/10.3390/diagnostics13193135>.
- Gelman, A., A. Jakulin, M. G. Pittau, and Y. S. Su. 2008. “A Weakly Informative Default Prior Distribution for Logistic and Other Regression Models.” *Annals of Applied Statistics* 2 (4): 1360–83. <https://doi.org/10.1214/08-AOAS191>.
- Hoeting, J. A., D. Madigan, A. E. Raftery, and C. T. Volinsky. 1999. “Bayesian Model Averaging: A Tutorial.” *Statistical Science* 14 (4): 382–417. <https://doi.org/10.1214/ss/1009212519>.
- Klauenberg, K., G. Wübbeler, B. Mickan, P. Harris, and C. Elster. 2015. “A Tutorial on Bayesian Normal Linear Regression.” *Metrologia* 52 (6): 878–92. <https://doi.org/10.1088/0026-1394/52/6/878>.
- Kruschke, J. K., and T. M. Liddell. 2017. “Bayesian Data Analysis for Newcomers.” *Psychonomic Bulletin & Review* 25 (1): 155–77. <https://doi.org/10.3758/s13423-017-1272-1>.
- Leeuw, C. de, and I. Klugkist. 2012. “Augmenting Data with Published Results in Bayesian Linear Regression.” *Multivariate Behavioral Research* 47 (3): 369–91. <https://doi.org/10.1080/00273171.2012.673957>.
- Liu, Y. M., S. L. S. Chen, A. M. F. Yen, and H. H. Chen. 2013. “Individual Risk Prediction Model for Incident Cardiovascular Disease: A Bayesian Clinical Reasoning Approach.” *International Journal of Cardiology* 167 (5): 2008–12. <https://doi.org/10.1016/j.ijcard.2012.05.016>.
- Luo, C., X. Sun, Y. Zhao, and H. Guo. 2024. “Variable Selection and Model Averaging in Bayesian Additive Regression Trees: A Comparative Study.” *Journal of Computational and Graphical Statistics*. <https://doi.org/10.1080/10618600.2024.2401234>.
- Momeni, F., F. Momeni, A. Moradi, M. Shabani, and B. Amani. 2021. “Bayesian State-Space Modeling of Dynamic COVID-19 Risk Prediction Using Time-Varying Biomarkers.” *Scientific Reports* 11 (1): 20387. <https://doi.org/10.1038/s41598-021-99711-7>.
- National Center for Health Statistics (NCHS). 2014. “National Health and Nutrition Examination Survey (NHANES) 2013–2014 Data Documentation, Codebook, and Frequencies.” U.S. Department of Health; Human Services, Centers for Disease Control; Prevention. <https://www.cdc.gov/nchs/nhanes/continuousnhanes/overview.aspx?BeginYear=2013>.
- Richardson, Robert, and Brian Hartman. 2018. “Bayesian Nonparametric Regression Models for Modeling and Predicting Healthcare Claims.” *Insurance: Mathematics and Economics* 83: 1–8. <https://doi.org/10.1016/j.insmatheco.2018.06.002>.
- Schoot, Rens van de, David Kaplan, Jaap Denissen, Jens Asendorpf, Franz Neyer, and Marcel van Aken. 2013. “A Gentle Introduction to Bayesian Analysis: Applications to Developmental Research.” *European Journal of Developmental Psychology* 10 (6): 723–49. <https://doi.org/10.1080/17405629.2013.803373>.
- Vande Schoot, R., S. Depaoli, R. King, B. Kramer, K. Märtnens, M. G. Tadesse, M. Vannucci,

- et al. 2021. “Bayesian Statistics and Modelling.” *Nature Reviews Methods Primers* 1: 1–26. <https://doi.org/10.1038/s43586-020-00001-2>.
- Zeger, S. L., Z. Wu, Y. Coley, A. T. Fojo, B. Carter, K. O’Brien, P. Zandi, et al. 2020. “Using a Bayesian Approach to Predict Patients’ Health and Response to Treatment.” 272. Johns Hopkins Biostatistics Working Paper Series. <https://biostats.bepress.com/jhubiostat/paper272>.